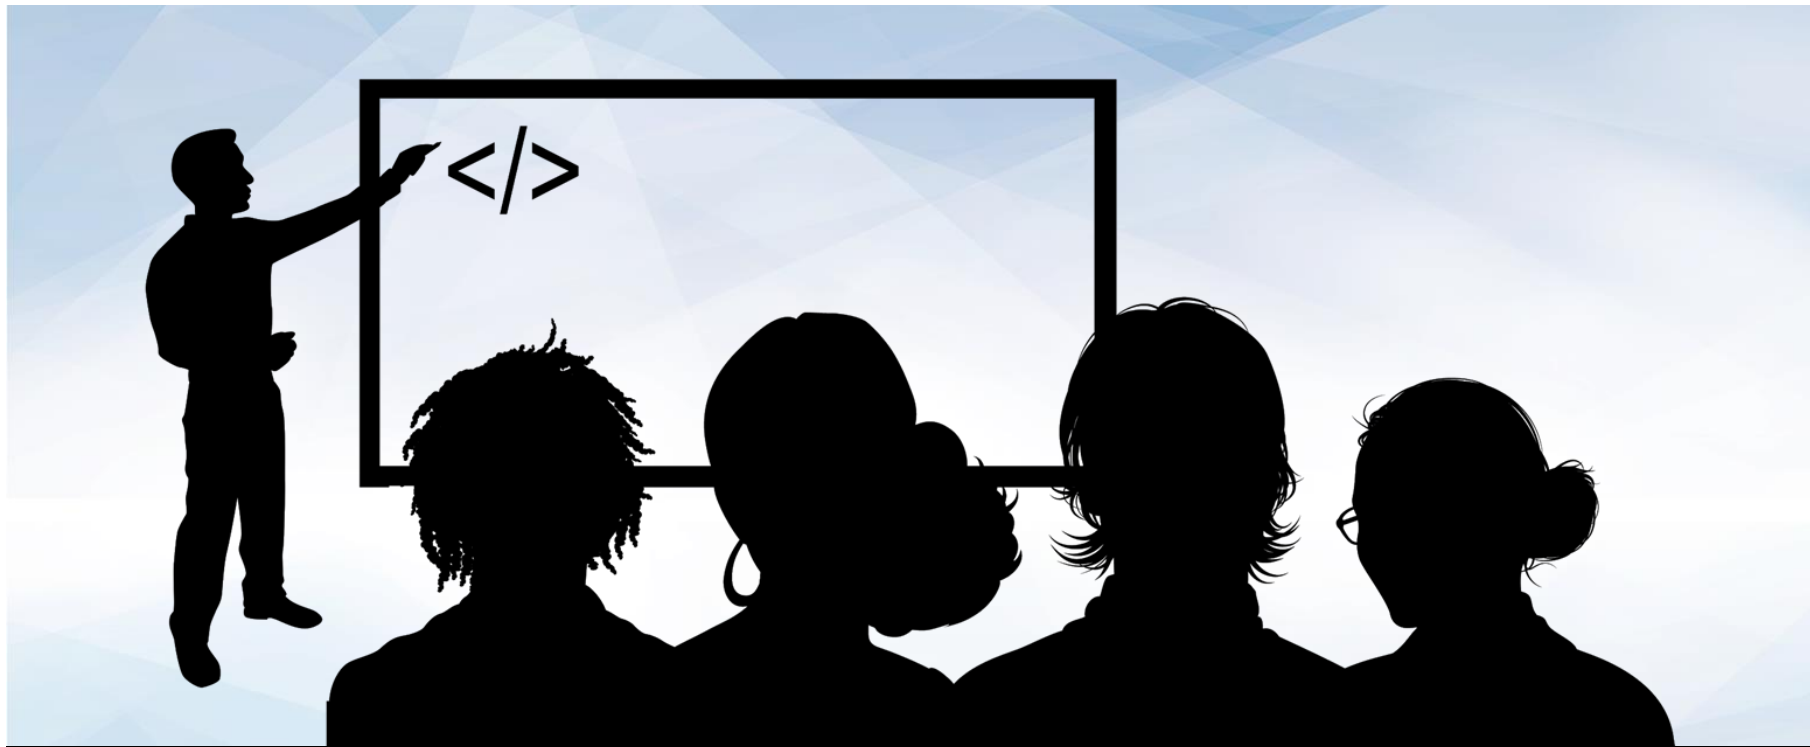# Introduction to Machine Learning

Instructor Demonstration
Introduction to Machine Learning

So It Begins...

Basic Definitions

Intelligent Algorithms (Definition)

" Intelligent algorithms are ones that use data to modify its behavior. Intelligent algorithms differ in that they can change their behavior as they run, often resulting in a user experience that many would say is intelligent. "

—*Algorithms of the Intelligent Web, Second Edition*

Algorithms of the Intelligent Web, Second Edition
by *Douglas G. McIlwraith Haralambos Marmanis Dmitry Babenko*
Publisher: Manning Publications
Release Date: August 2016

Intelligent Algorithms (Diagram)

Intelligent algorithms are ones that respond to data such that the algorithm gets better. It effectively "evolves." The decision is no longer deterministic given the event.

Data

Event

Intelligent ▮ Algorithm

Decision

Intelligent Algorithms (Triad)

Machine Learning
Capability of software to generalize phenomena (past or future) based on past experience

Predictive Analytics
Capability of software to predict future outcomes based on historical data

Artificial Intelligence
Software (and machines) that have a series of options to achieve a particular goal
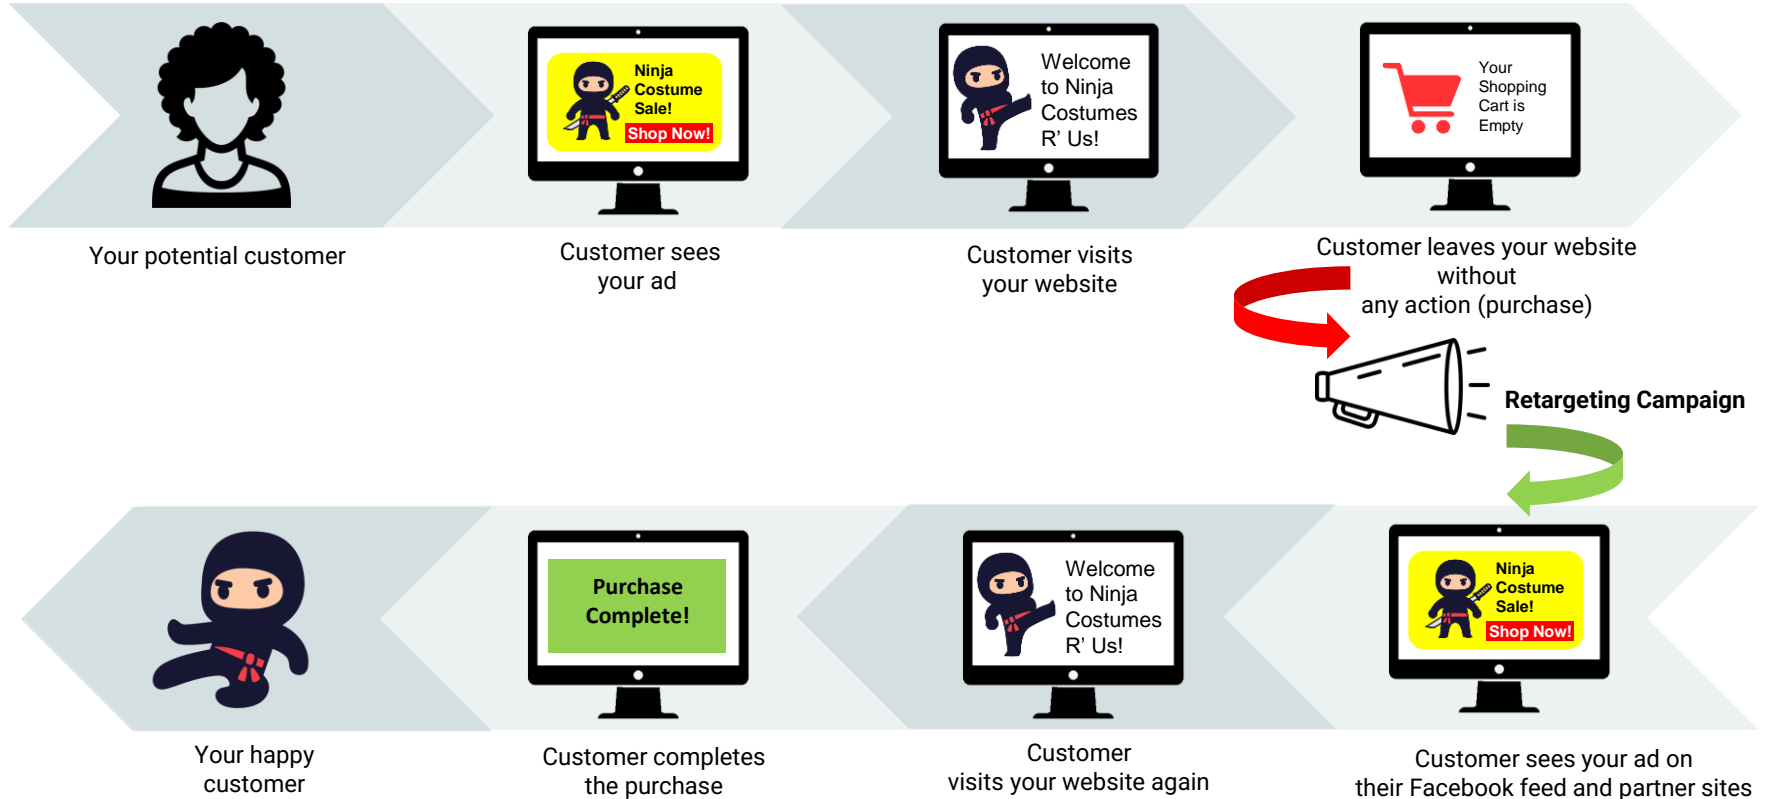
Artificial Intelligence (Example)

Predictive Analytics (Example)

How retargeting ads work:

Ninja Costume Sale! Shop Now!

Welcome to Ninja Costumes R' Us!

Your Shopping Cart is Empty

Your potential customer

Customer sees your ad

Customer visits your website

Customer leaves your website without any action (purchase)

Retargeting Campaign

Purchase Complete!

Welcome to Ninja Costumes R' Us!

Ninja Costume Sale! Shop Now!

Your happy customer

Customer completes the purchase

Customer visits your website again

Customer sees your ad on their Facebook feed and partner sites

Machine Learning (Categories)

Dimensionality Reduction
Meaningful Compression
Big Data Visualization
Structure Discovery
Feature Elicitation

Unsupervised Learning

Clustering
Recommender Systems
Targeted Marketing
Customer Segmentation

Machine Learning

Reinforcement Learning
Real-time decisions
Robot Navigation
Game AI
Skill Acquisition

Classification
Image Classification
Identity Fraud Detection
Customer Retention
Diagnostics

Supervised Learning

Regression
Population Growth Prediction
Advertising Popularity Prediction
Weather Forecasting
Market Forecasting
Estimating Life Expectancy

Machine Learning (Categories)

Dimensionality Reduction

Meaningful Compression
Big Data Visualization
Structure Discovery
Feature Elicitation

Classification:
Applying labels to data

Image Classification
Identity Fraud Detection
Customer Retention
Diagnostics

Unsupervised Learning

Machine Learning

Supervised Learning

Clustering:
Finding groups within a population

Recommender Systems
Targeted Marketing
Customer Segmentation

Regression:
Fitting data to predict where a new data point lies

Population Growth Prediction
Advertising Popularity Prediction
Weather Forecasting
Market Forecasting
Estimating Life Expectancy

Reinforcement Learning

Real-time decisions
Robot Navigation
Game AI
Skill Acquisition

Machine Learning (Supervised)

Supervised Learning: Algorithms for which the potential outcomes are knowable in advance (i.e., category or numeric range) and can be used to correct the model's predictions.

01 Example

Using data such as credit score, credit history, income, etc., we are trying to predict whether an individual is a credit risk or not.

Known Category:
"Credit Risk" vs. "Not Credit Risk"

02 Example

Using features such as number of bedrooms, square feet, etc., we are trying to predict the market value of a house.

Numeric Range:
50,000–500,000

Machine Learning (Regression)

We'll be revisiting regression to predict the location of data points based on old data.

Machine Learning (Classification)

Machine Learning (Classification)

In classification problems, our focus is identifying which predefined label our data falls into, based on the **features** we have.

From: cheapsales@buystufffromme.com To: ang@cs.stanford.edu Subject: Buy now!	From: Alfred Ng To: ang@cs.stanford.edu Subject: Christmas dates?
Deal of the week! Buy now! Rolex w4tchs - $100 Med1cine (any kind) - $50 Also low cost M0rgages available.	Hey Andrew, Was talking to Mom about plans for Xmas. When do you get off work? Meet Dec 22? -Alf
Spam	Non-Spam

Machine Learning (Unsupervised)

Unsupervised Learning: Algorithms for which the potential outcomes are unlabeled. Inferences are made directly from the data without feedback from known outcomes or labels.

Machine Learning (Clustering)

In this clustering problem, we expect our algorithm to find the groupings of data points based on location.

Machine Learning (Clustering)

In this clustering problem, we expect our algorithm to find the groupings of data points based on location.

K=3

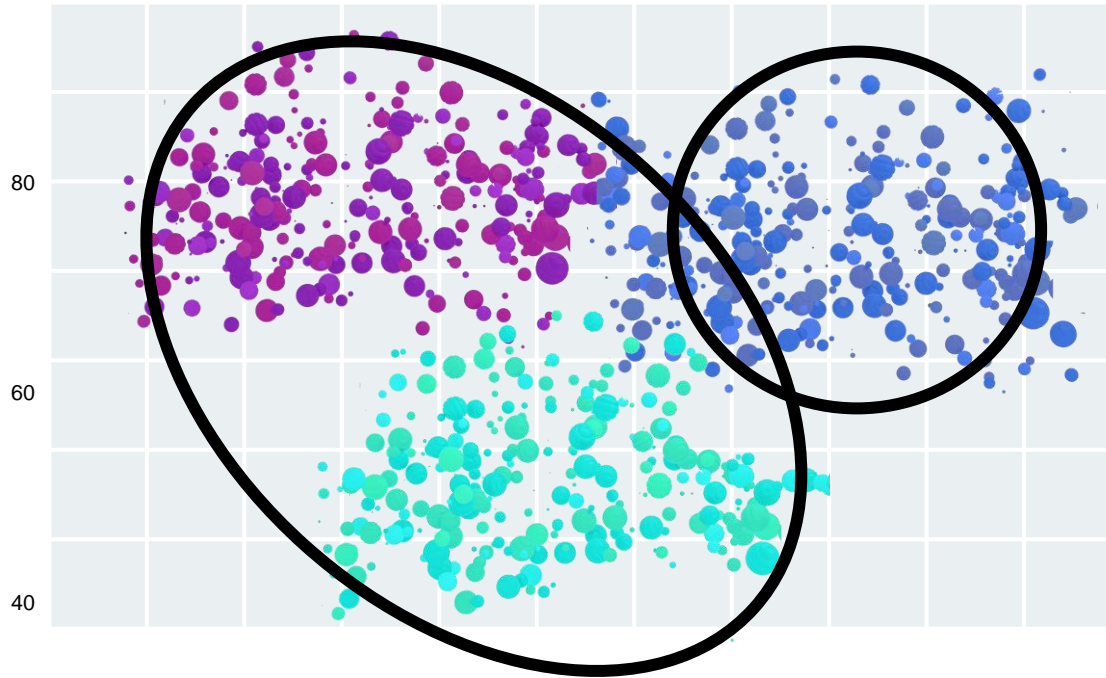# Machine Learning (Clustering)

But the problem is more complex:

K=2

Machine Learning (Clustering)

Perhaps the clusters are not where we think they are.

K=2

Training and Predicting

Regardless of the problem type, in Machine Learning we follow a familiar paradigm.

Model → Fit (Train) → Predict

Machine Learning

Data → Fit Predict Model → Query → Decision → Application

Training and Predicting

Regardless of the problem type, in Machine Learning we follow a familiar paradigm.

Model → Fit (Train) → Predict

A	B	C	Class
11	16	22	1
10	8	4	2
…	…	…	…

A	B	C	Class
10	15	23	?

Apply

Fit

Model

Predict

Class
1

Many Models: Which Do We Choose?

Dimensionality Reduction
Meaningful Compression
Big Data Visualization
Structure Discovery
Feature Elicitation

Classification
Image Classification
Identity Fraud Detection
Customer Retention
Diagnostics

Unsupervised Learning

Machine Learning

Supervised Learning

Clustering
Recommender Systems
Targeted Marketing
Customer Segmentation

Regression
Population Growth Prediction
Advertising Popularity Prediction
Weather Forecasting
Market Forecasting
Estimating Life Expectancy

Reinforcement Learning
Real-time decisions
Robot Navigation
Game AI
Skill Acquisition

Machine Learning will one day likely be relegated to a world of the off-the-shelf formulas readily available by the masses. We're already nearly there.

Questions?

Instructor Demonstration

Univariate Linear Regression

We start with something familiar...

...linear regression!

What is **linear regression**?

Linear Regression is Used to Model and Predict A Relationship

Linear regression…

- Predicts a dependent variable, given values from an independent variable.

- There are two basic types
  - Simple linear regression
  - Multiple linear regression

- Both types predict an independent variable using the linear equation

The Equation of a Line

$$y = mx + b$$

Dependent variable

Slope

Independent variable

y-intercept

Linear regression is *FAST*

<Time to Code>

Instructor Demonstration

Quantifying Regression

Common Scoring Metrics

R^2 (R-Squared):

This is the baseline metric that many ML tools report on score. Higher R^2 values signify that the model is "highly predictive." An R^2 value of >0.90 means that our model roughly accounts for 90% of the variability of the data.

02

MSE (Mean Squared Error):

This measures the average of the squares of the errors or deviations.

Basic Premise of Validation Using Training/Testing Data

We will cut a slice of this data (80%) to build our model, and then use this slice to predict the values for the remaining 20%.

Full Data Set (Historic)

N=1000

# bedrooms	# baths	Sq. feet (k)	Price (k)
2	1	1	200
3	2	1.5	250
…	…	…	…

Training Data Set

N=800

# bedrooms	# baths	Sq. feet (k)	Price (k)
4	3.5	3.2	450
2	2	1.5	220
…	…	…	…

Testing Data Set

N=200

| # bedrooms	# baths	Sq. feet (k)	Price (k)
1	1	.5	60
5	3.5	4.2	780
…	…	…	…

Fit

Model

Predict

Basic Premise of Validation: Training

We use the training data to fit the model to the data. This is the training step where we build a model that can predict our output (home price) for a given set of features (# bedrooms, # baths, square feet). Once the model is trained, we can use the model to make predictions.

Full Data Set (Historic)

N=1000

# bedrooms	# baths	Sq. feet (k)	Price (k)
2	1	1	200
3	2	1.5	250
...	...	...	...

Training Data Set

N=800

# bedrooms	# baths	Sq. feet (k)	Price (k)
4	3.5	3.2	450
2	2	1.5	220
...	...	...	...

Testing Data Set

N=200

# bedrooms	# baths	Sq. feet (k)	Price (k)
1	1	.5	60
5	3.5	4.2	780
...	...	...	...

Fit → Model

Predict

Basic Premise of Validation

We use the test data to make new home price predictions. We can then compare the home price of our prediction vs. the actual price. Based roughly on how often we are "correct," we get a score for the model as a whole. If the model scores well, we can trust it for future use. We train the model on the training data and score the model based on data that it has never seen before (test data).

Full Data Set (Historic)

N=1000

# bedrooms	# baths	Sq. feet (k)	Price (k)
2	1	1	200
3	2	1.5	250
…	…	…	…

Training Data Set

N=800

# bedrooms	# baths	Sq. feet (k)	Price (k)
4	3.5	3.2	450
2	2	1.5	220
…	…	…	…

Testing Data Set

N=200

# bedrooms	# baths	Sq. feet (k)	Price (k)
1	1	.5	60
5	3.5	4.2	780
…	…	…	…

Fit → Model

Predict